

GPU batched vs CPU sequential — 16x8, D=8, $\chi=32$, A100

